



DEEP
LEARNING
INSTITUTE




DLI Accelerated Data Science Teaching Kit

Lecture 12.1 - HBase Overview



The Accelerated Data Science Teaching Kit is licensed by NVIDIA, Georgia Institute of Technology, and Prairie View A&M University under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).

Lesson Based on



<http://shop.oreilly.com/product/0636920014348.do>



<http://hbase.apache.org>

Built on top of HDFS

Supports **real-time** read/write random access

Scale to very large datasets, many machines

Not relational, does NOT support SQL

(“**NoSQL**” = “not only SQL”) <http://en.wikipedia.org/wiki/NoSQL>

Supports **billions of rows, millions of columns**

(e.g., serving Facebook’s Messaging Platform)

Written in Java; works with other APIs/languages (REST, Thrift, Scala)

<http://radar.oreilly.com/2014/04/5-fun-facts-about-hbase-that-you-didnt-know.html>

<http://hbase.apache.org/poweredbyhbase.html>

HBase is Column-Oriented

Data is stored by **columns** physically (instead of by rows)

- A row conceptually consists on many columns (millions!)

Rows form a table

- **Row key** uniquely locates a row, like an “**index**”, **only one** “index” per table
 - No **built-in** support for multiple indices; possible via **extensions**
- Rows **sorted** by row key lexicographically (~ alphabetically)

Rows Sorted Lexicographically

(=alphabetically)

```
hbase(main):001:0> scan 'table1'
ROW      COLUMN+CELL
row-1    column=cf1:, timestamp=1297073325971 ...
row-10   column=cf1:, timestamp=1297073337383 ...
row-11   column=cf1:, timestamp=1297073340493 ...
row-2    column=cf1:, timestamp=1297073329851 ...
row-22   column=cf1:, timestamp=1297073344482 ...
row-3    column=cf1:, timestamp=1297073333504 ...
row-abc  column=cf1:, timestamp=1297073349875 ...
7 row(s) in 0.1100 seconds
```

“row-10” comes before “row-2”.

How to fix?

Pad “row-2” with a “0”.

i.e., “row-02”

Columns grouped into **column families**

- Why?
 - Helps with organization, understanding, optimization, etc.
- In details...
 - Columns in the same family stored in same *file* called *HFile*
 - Apply compression on the whole family

More on column family, column

Column family

- An HBase table supports only **few** families (e.g., <10)
 - Due to limitations in implementation
- Family name must be **printable**
- Should be defined when table is created
 - Should not be changed often

Each **column** referenced as “**family:qualifier**”

- Can have **millions** of columns

Cell Value

Timestamped

- Implicitly by system
- Or set explicitly by user

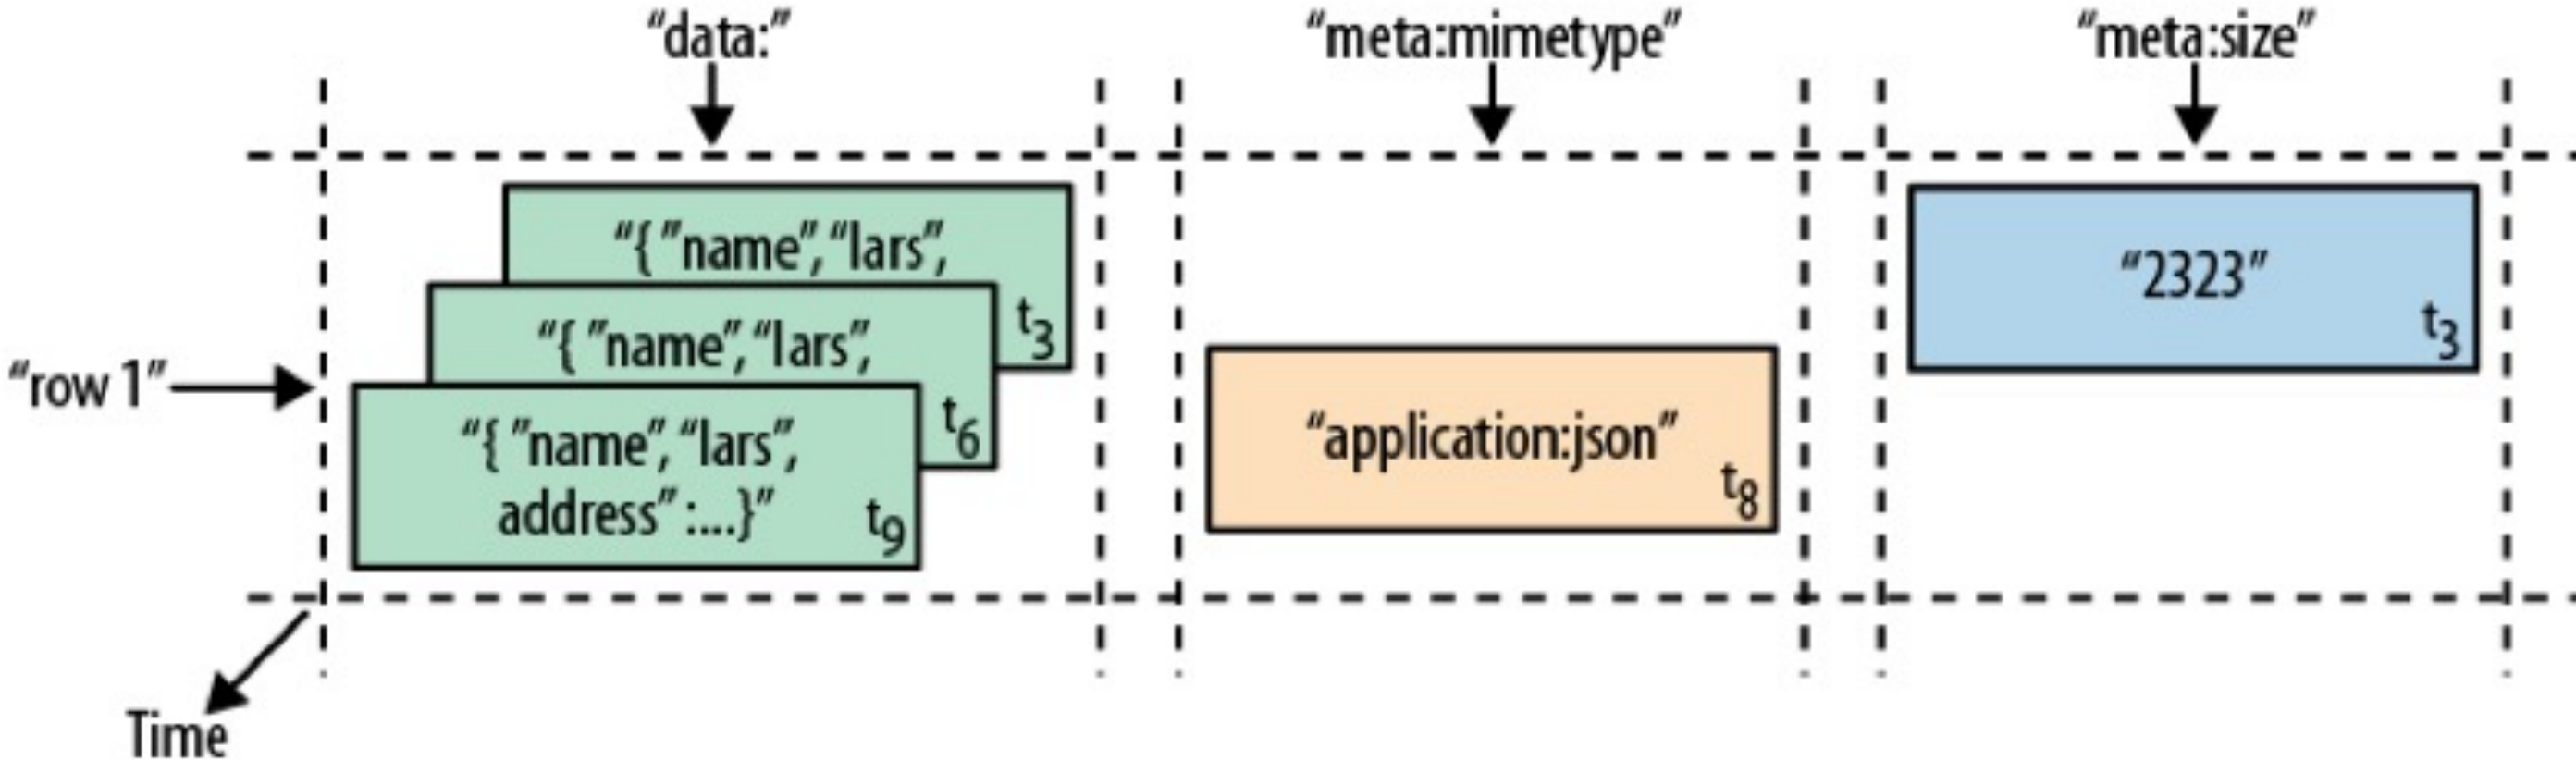
Let you store **multiple versions** of a value

- = values over time

Values stored in **decreasing** time order

- **Most recent value** can be read first

Time-oriented view of a row



Row Key	Time Stamp	Column "data:"	Column "meta:"	
			"mimetype"	"size"
"row1"	t_3	<code>{ "name": "lars", "address": ... }</code>		"2323"
	t_6	<code>{ "name": "lars", "address": ... }</code>		
	t_8		"application/json"	
	t_9	<code>{ "name": "lars", "address": ... }</code>		

An Exercise

How would you use HBase to create a webtable to store snapshots of every webpage on the planet, over time?

Row key	Time stamp	html	javascript	image1	image2	...
"com.cnn.www"	t8	<html>...	<javascript>...			...
						...
						...
						...

An Exercise

How would you use HBase to create a webtable to store snapshots of every webpage on the planet, over time?



DEEP
LEARNING
INSTITUTE



DLI Accelerated Data Science Teaching Kit

Thank You