



DEEP
LEARNING
INSTITUTE



DLI Accelerated Data Science Teaching Kt

Lecture 14.4 - RAPIDS Acceleration: Linear Regression



The Accelerated Data Science Teaching Kit is licensed by NVIDIA, Georgia Institute of Technology, and Prairie View A&M University under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).

RAPIDS

The RAPIDS data science framework includes a collection of libraries for executing end-to-end data science pipelines completely in the GPU.

It is designed to have a familiar look and feel to data scientists working in Python.



Features

Hassle-Free Integration Accelerate your Python data science toolchain with minimal code changes and no new tools to learn.	Top Model Accuracy Increase machine learning model accuracy by iterating on models faster and deploying them more frequently.
Reduced Training Time Drastically improve your productivity with near-interactive data science.	Open Source Customizable, extensible, interoperable - the open-source software is supported by NVIDIA and built on Apache Arrow.

Speed Up Learning of Linear Regression

Linear Regression is a simple machine learning model where the response y is modelled by a linear combination of the predictors in X .

The model can take array-like objects, either in host as NumPy arrays or in device (as Numba or `cuda_array_interface`-compliant), as well as cuDF DataFrames as the input.

Linear Regression vs Linear Regression cuML

Import packages

```
import cudf
from cuml import make_regression, train_test_split
from cuml.linear_model import LinearRegression as cuLinearRegression
from cuml.metrics.regression import r2_score
from sklearn.linear_model import LinearRegression as skLinearRegression
```

Setting parameters

```
n_samples = 2**20 #If you are running on a GPU with less than 16GB RAM, please change to 2**19 or you could run out of memory
n_features = 399

random_state = 23
```

Linear Regression vs Linear Regression cuML

Generating Data

```
%%time
X, y = make_regression(n_samples=n_samples, n_features=n_features, random_state=random_state)

X = cudf.DataFrame(X)
y = cudf.DataFrame(y)[0]

X_cudf, X_cudf_test, y_cudf, y_cudf_test = train_test_split(X, y, test_size = 0.2, random_state=random_state)

# Copy dataset from GPU memory to host memory.
# This is done to later compare CPU and GPU results.
X_train = X_cudf.to_pandas()
X_test = X_cudf_test.to_pandas()
y_train = y_cudf.to_pandas()
y_test = y_cudf_test.to_pandas()
```

Linear Regression vs Linear Regression cuML

Sklearn Linear Regression

Fit

```
%%time
ols_sk = sklearn.LinearRegression(fit_intercept=True,
                                  normalize=True,
                                  n_jobs=-1)

ols_sk.fit(X_train, y_train)
```

CPU times: user 29 s, sys: 5.47 s, total: 34.5 s
Wall time: 21.3 s

Predict

```
%%time
predict_sk = ols_sk.predict(X_test)
```

CPU times: user 125 ms, sys: 490 μ s, total: 125 ms
Wall time: 105 ms

Evaluate

```
%%time
r2_score_sk = r2_score(y_cudf_test, predict_sk)
```

CPU times: user 5.18 ms, sys: 20 μ s, total: 5.2 ms
Wall time: 11.2 ms

cuML

Fit

```
%%time
ols_cuml = cuML.LinearRegression(fit_intercept=True,
                                  normalize=True,
                                  algorithm='eig')

ols_cuml.fit(X_cudf, y_cudf)
```

CPU times: user 181 ms, sys: 161 ms, total: 343 ms
Wall time: 454 ms

Predict

```
%%time
predict_cuml = ols_cuml.predict(X_cudf_test)
```

CPU times: user 29.2 ms, sys: 6.27 ms, total: 35.5 ms
Wall time: 36.4 ms

Evaluate

```
%%time
r2_score_cuml = r2_score(y_cudf_test, predict_cuml)
```

CPU times: user 1.25 ms, sys: 1.07 ms, total: 2.32 ms
Wall time: 2.33 ms

Linear Regression vs Linear Regression cuML

Compare Results

```
print("R^2 score (SKL): %s" % r2_score_sk)  
print("R^2 score (cuML): %s" % r2_score_cuml)
```

```
R^2 score (SKL): 1.0  
R^2 score (cuML): 1.0
```




DEEP
LEARNING
INSTITUTE



DLI Accelerated Data Science Teaching Kit

Thank You